

# Improving application responsiveness and I/O latency with the BFQ I/O scheduler

---

Paolo Valente

Department of Physics, Computer Science  
and Mathematics  
Modena - Italy

**UNIMORE**  
UNIVERSITÀ DEGLI STUDI DI  
MODENA E REGGIO EMILIA




 **P-SOCRATES**

LinuxDay Pisa 2014

# Budget Fair Queueing (BFQ) 1/2

---

- Storage-I/O scheduler
  - High responsiveness
  - Low latency for soft real-time (time-sensitive) applications, such as multimedia ones
  - High throughput
  - Desired throughput fraction guaranteed to each application
    - even if throughput fluctuates



# Budget Fair Queueing (BFQ) 2/2

---

- Adopted in a number of distributions and kernel variants
- Submitted to *lkml* about four months ago:  
<https://lkml.org/lkml/2014/5/27/314>
  - Arranged a roadmap for possible inclusion
  - By replacing CFQ
- BFQ homepage:  
[http://algogroup.unimore.it/people/paolo/disk\\_sched/](http://algogroup.unimore.it/people/paolo/disk_sched/)

# Contents of this presentation

---

1. Two demos of the performance of BFQ
  - Compared with CFQ, DEADLINE and NOOP
  - On an SSD and on an HDD
2. Some considerations about BFQ, fast devices and latency

- Links to the videos of the demos in BFQ homepage

[http://algogroup.unimore.it/people/paolo/disk\\_sched/](http://algogroup.unimore.it/people/paolo/disk_sched/)

# The trick ...

---

- Applications are launched quickly, and interactive and soft real-time applications enjoy a low latency because
  - BFQ privileges the I/O related to interactive or soft real-time tasks
- Hard part
  - Not losing throughput
  - Correctly detecting applications to privilege
  - Implementing all the logic cleanly

# Speed and latency

---

- Because of its execution time, the current version of BFQ is likely to be a bottleneck on high-end, high-speed devices
  - “But little or no scheduling is needed on such devices”
    - “In fact, as the speed increases, latency problems will just go away”
  - True?

# Speed and latency

---

- Not that sure
  - Device-related issues
  - Workload-related issues



# Device-related key problem

---

- Devices usually **reorder** I/O requests
- Even if a device is very fast, but it
  - systematically serves many wrong requests
  - before serving the right ones,then responsiveness and latency for soft real-time applications are likely to be still bad
  - Exactly the cause of the problems shown in the demo
- Ordering might be controlled by passing in priorities
  - But this would hurt performance

# Relation with new devices

---

- Speed will increase
- But expectedly through higher parallelism
  - Devices will be fed with more requests
  - Internal device schedulers may then happen to serve more wrong requests before the right ones
  - The *wrong-service-order* problem may remain unaltered
    - Or even get worse

# Workload-related issues 1/2

---

- High-speed, costly devices make sense where high throughput is needed
- For example, where many instances of the same application need to be executed in parallel
  - Virtual machines in clouds
  - Instances of streaming servers in Video-on-Demand services

# Workload-related issues 2/2

---

- For these applications
  - If the available throughput becomes  $N$  times as high as before
    - Also the number of instances that can be executed becomes  $N$  times as high
    - The per-instance throughput, and hence the request-completion latencies would then be about the same as before
  - In the end, latency issues are likely to remain about the same as before
    - Or may become even worse, because there would be more outstanding I/O requests

# Future work on BFQ

---

- Dealing with millions of IOPS
  - Measuring the impact of BFQ
  - Investigating simpler variants of BFQ
    - Useful also if one may want to use BFQ as an internal scheduler in a device
      - This could enable low-latency guarantees to be provided with no or a negligible throughput penalty
- Guaranteeing high responsiveness and low latency also in virtualized environments

# The end

---

Thanks for your attention

Questions?